

CS 372 Lecture #22

Reliable data transfer with TCP

- congestion control

Note: Many of the lecture slides are based on presentations that accompany *Computer Networking: A Top Down Approach*, 6th edition, by Jim Kurose & Keith Ross, Addison-Wesley, 2013.

Congestion Control

- Goals:
 - Optimize utilization
 - Handle causes of network congestion
 - Retransmission handles only the symptoms
 - Detect congestion
 - Avoid congestion (when possible)
 - React to congestion (when necessary)

Detecting congestion

- Two approaches
 - Inferred by end-to-end systems
 - delay/loss
 - used by TCP
 - Network core assistance
 - routers provide feedback to senders
 - more later about ICMP
 - Internet Control Messaging Protocol

Avoiding congestion


- Handle the causes ... and optimize utilization
 - timeouts (see previous lecture)
 - sliding windows
 - If receiving application drains a few data bytes, receiver will advertise a small window , and sender will immediately send small segment to fill window
 - Inefficient in processing time and network bandwidth
 - TCP solution:
 - Receiver delays advertising new window
 - Sender delays sending data when window is small

Reacting to (handling) congestion

- Delicate balance
 - **Too slow**: under-utilization
 - waste of network resources
 - **Too fast**: over-utilization
 - waste of network resources
- Approach: **Adaptive transmission algorithm**
 - Define **CongWin** (sender's window size)
 - Slowly increase **CongWin** to probe for usable bandwidth
 - Decrease the sending rates when congestion is detected
 - TCP uses additive-increase, multiplicative-decrease (**AIMD**)
- Definition: **MSS (Maximum Segment Size)**
 - Upper limit on amount of data that can be sent in the largest link-layer frame on the sending host (usually 1460 bytes)

TCP congestion control: AIMD


- Additive-increase, multiplicative decrease (**AIMD**)
 - **additive increase**: increase **CongWin** by 1 MSS every RTT until loss is detected
 - **multiplicative decrease**: cut **CongWin** in half after loss



Saw tooth
behavior: probing
for bandwidth

TCP congestion control: Slow Start

- When connection begins
 - **CongWin** = 1 MSS
 - available bandwidth may be much greater than MSS/RTT
 - avoids congestion, but need to quickly increase to a respectable rate
- increase rate exponentially until first loss event
 - double **CongWin** every RTT
 - must stop increasing rate if congestion is detected
- initial rate is slow but increases exponentially fast
 - must define a **Limit**
 - **CongWin** grows linearly after **Limit**



Refinement: inferring loss

- After 3 dup ACKs:
 - **CongWin** is cut in half
 - window then grows linearly
- But after timeout event:
 - **CongWin** is set to 1 MSS
 - window then grows exponentially to a **Limit**, then grows linearly

Philosophy:

- 3 duplicate ACKs implies that the network is capable of delivering some segments, so OK to assume packet is lost.
- timeout indicates a “more alarming” congestion scenario

Summary: TCP Congestion Control

- When **CongWin** is below **Limit**, sender is in **slow-start** phase
 - **CongWin** grows exponentially.
- When **CongWin** is above **Limit**, sender is in **congestion-avoidance** phase
 - **CongWin** grows linearly.
- When a **triple duplicate ACK** occurs, **Limit** set to **CongWin/2** and **CongWin** set to **Limit**.
 - back to slow-start or congestion-avoidance
- When **timeout** occurs, **Limit** set to **CongWin/2** and **CongWin** is set to 1 MSS.
 - back to slow-start
- Several TCP versions (RENO, TAHOE, VEGAS, etc.) See http://en.wikipedia.org/wiki/TCP_congestion_avoidance_algorithm

- Definitions:
 - MSS, AIMD, congestion window
- Congestion control
 - detection
 - delay/loss
 - avoidance
 - smarter timers, sliding-windows, slow-start
 - reaction
 - adaptive transmission