

IP Datagram Fragmentation

- Maximum Transmission Unit (MTU)
- Datagram re-assembly

Note: Many of the lecture slides are based on presentations that accompany *Computer Networking: A Top Down Approach*, 6th edition, by Jim Kurose & Keith Ross, Addison-Wesley, 2013.

Encapsulation across multiple hops

- Whenever a datagram transits a physical network, it is encapsulated in a frame appropriate to that network.
 - Hardware frame encapsulation/un-encapsulation at the link layer
 - Hardware frame formats may differ
- Each router in the path from the source to the destination:
 - *Un-encapsulates* incoming datagram from frame (link layer)
 - *Processes* datagram - determines next hop (network layer)
 - *Re-encapsulates* datagram in outgoing frame (link layer)
- Datagram itself is (almost!) unchanged
 - may be fragmented

- Every hardware technology specification includes the definition of the maximum size of the frame data area
 - *maximum transmission unit* (MTU)
- Any datagram encapsulated in a hardware frame must be smaller than the MTU for that hardware
- Internet has variety of network technologies with different MTUs
 - Source doesn't know *path MTU*
 - *smallest MTU for complete route* to destination
 - Easy for source to limit IP datagram size to *local MTU*
 - link layer must pass local MTU up to transport layer for TCP segments
 - A downstream network might have smaller MTU than local network

Choosing datagram size

- IP datagrams can be much larger than hardware MTUs
 - IP **TOTAL LENGTH**: maximum is 65,535 bytes (including header)
- Typical hardware frames:
 - **Ethernet**: 1500 bytes
 - **Token ring**: 2048 or 4096 bytes
- One technique - limit datagram size to smallest MTU of *any* known network
 - Problem: impedes use of newer technologies
- IPv4 uses *fragmentation*

Fragmentation

- Router detects inbound datagram that is larger than the outbound *network MTU*
 - Splits datagram into fragments
 - Each fragment is smaller than the MTU of the outbound network
- Each fragment is an independent datagram
 - Includes all header fields
 - All fragments have the **IDENTIFIER** of the original datagram
 - 3-bit **FLAGS** field indicates if datagram is a fragment
 - 13-bit **FRAGMENT OFFSET** gives order of fragment in original datagram

IP datagram format

ver	head. len	service type	length	
16-bit identifier			flgs	fragment offset
time to live		upper layer	header checksum	
32 bit source IP address				
32 bit destination IP address				
Options (if any)				
data (variable length, typically a TCP or UDP segment)				

IP Fragmentation and Reassembly

Example

- 4000-byte datagram
- MTU = 1500 bytes

	length	ID	fragflag	offset	
	=4000	=x	=0	=0	

One large datagram becomes several smaller datagrams

1480 bytes in
data field

offset in
8-byte blocks
= 1480/8

	length	ID	fragflag	offset	
	=1500	=x	=1	=0	

	length	ID	fragflag	offset	
	=1500	=x	=1	=185	

	length	ID	morefrag	offset	
	=1060	=x	=0	=370	

fragflag: 0 = not a fragment, 1 = is a fragment
dontfrag: 0 = OK to fragment, 1 = do not fragment
morefrag: 0 = last fragment, 1 = more fragments

Re-assembly

- Reconstruction of original datagram is called *re-assembly*
 - performed by protocol stack at final destination
- **IDENTIFICATION** field in each fragment matches original datagram
 - fragments from different datagrams can arrive out of order and still be sorted out
- **FLAGS** field bit identifies fragment containing end of data from original datagram
- **FRAGMENT OFFSET** specifies re-assembly order
 - counted in 8-byte blocks

Re-assembly

- Once a datagram is fragmented, it stays fragmented until it reaches destination.
- The final destination performs the reassembly.
 - Reassembly is not done along the way, even if an outbound network has a larger MTU.
- Two advantages to having the final destination do the reassembly:
 - It reduces the amount of information needed by routers.
 - It allows routers to change dynamically.
 - It avoids re-fragmenting

Fragment loss

- IP does not guarantee datagram delivery, so fragments may be lost.
- **Destination**
 - Sets timer with first fragment
 - If timer expires before all fragments arrive
 - one or more fragments assumed lost
 - Destination drops *entire* original datagram
 - **Source** is assumed to retransmit entire original datagram
- Reasons for the all-or-nothing behavior:
 - There's no mechanism for a receiver to tell a sender which fragments arrived.
 - The sender does not know about any fragmentation that happened along the way.
 - Even if the sender could retransmit a single fragment, it might take a different route and be fragmented differently.

Fragmenting a fragment

- Fragment may encounter subsequent network with even smaller MTU
- Router fragments the fragment to fit local MTU
- Resulting (sub)fragments look just like original datagrams (except for size and offsets)
- No need to reassemble hierarchically
- IP does not distinguish among "original" fragmentation and subsequent fragmentations
- Re-assembly is entirely the responsibility of the protocol on the destination host

- Definitions:
 - MTU
 - local MTU, network MTU, path MTU
- fragmentation, fragments, sub-fragments
- re-assembly